

Effective Logging for Shiny

Tan Ho

ML Engineer, Zelus Analytics 

Shiny in Production 2023


tanho.ca/logging-shiny




Motivations

BUG FIXING WAYS


LOG FILES




BREAKPOINTS



PAIR PROGRAMMING




STACKOVERFLOW



PACT WITH THE DEVIL



REBRANDING



MONKEYUSER.COM

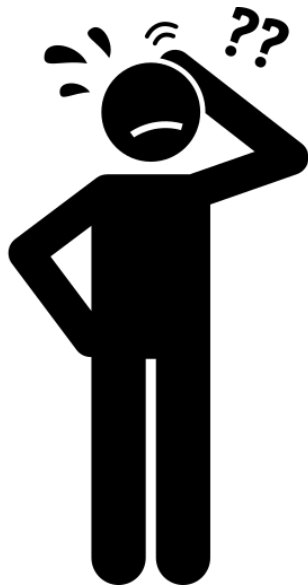
<https://www.monkeyuser.com/2019/bug-fixing-ways/>



Me, trying to find the log files in our production Shiny apps



Me, finding out the log files weren't useful



Me, trying to find best practices for logging in R & Shiny

WRITE THE
BOOK YOU
WANT TO READ.

- Austin Kleon, *Steal Like An Artist*

Agenda



App-level logging



Production-level logging

App-level logging




**Who are you
logging for?**

Humans

- Easy to scan visually
- Easy to understand
- Likes good summaries

Printing (to console)



Machines

- Easy to ingest
- Accuracy
- Completeness

Structured Data (JSON)

Who are you logging for?

**What should I be
logging in my app?**

**What questions
should my logs
help me answer?**

Errors

- Did we handle the error gracefully?
- Do we know the cause of the error?
- Can we get the context of this error?
- Can we reproduce this error?

Versioning and Environment Issues

Is this app running with...

- the latest versions of internal packages?
- the latest (or correct) versions of dependencies?
- the correct environment variables and package options?


Users

- Who is using the app?
- What pages/sections/processes are they using?
- What actions did they take?
- Are they doing anything unexpected or suspicious?

Expensive Functions & External Requests

- Which API/database/background job are we accessing?
- What (arguments, query) did we request from it?
- How long did it take? When did it start? When did it end?
- Did each step of the process complete successfully?

Problem: alert
fatigue



Solution: alert levels!



Debug



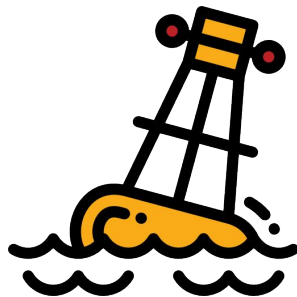
Devel




Info



Production



Warning



Error



Fatal

What goes in each alert level?




Debug



Info



Warn



Error



Fatal

- Print debugging
 - input values, reactive states
 - helpful context for diagnostics
- Progress tracking within larger functions

What goes in each alert level?




Debug



Info



Warn



Error



Fatal

- Session-level information (eg user auth, system info)
- Performance timings of major/expensive calculations
- Significant user actions


What goes in each alert level?




Debug



Info



Warn



Error



Fatal

- Kind of like R's *warning()*
- Edge cases, suspicious behaviours, and unusual inputs/states that still produce output


What goes in each alert level?



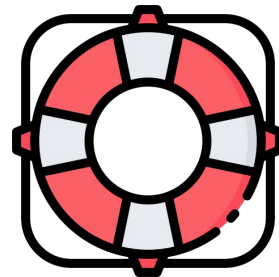
Debug



Info



Warn



Error



Fatal

- Errors that are handled gracefully by the app, e.g.
 - Input/argument validation
 - Function errors wrapped in *tryCatch()* handlers

What goes in each alert level?



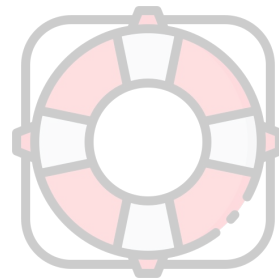
Debug



Info



Warn



Error



Fatal

- Errors causing app crashes (ie Shiny grey screens)
- Corresponds to R's *stop()* calls (unless properly handled)

**How should I log
things in my app?**

Choosing a logging package



- futile.logger
- log4r
- logger
- logging
- loggit
- lumberjack
- rlog
- rsyslog

Choosing a logging package



logger

- ★ familiar syntax (glue)
- ★ performant & lightweight
- ★ actively maintained
- ★ flexible
 - alert levels
 - custom layouts & formatting
 - multiple destinations
 - can log as an async process

Some useful tricks

sitrep(): a customized session_info()

System info

Env variables

Internal packages

Package options

Dependency versions

```
— System Info —
• R version 4.2.1 (2022-06-23) • Running under: Ubuntu 20.04.5 LTS
— Environment Variables —
• ZELUS_ENV : WEBSITE
• ZELUS_DB_HOST : soccer-website
• ZELUS_DB_LOCATION : us-west-2
• SHINYPROXY_OIDC_ACCESS_TOKEN: eyJ..<redacted>..zLQ
• SHINYPROXY_PUBLIC_PATH : /app_proxy/e5ad2e1a-71c7-4412-8c75-e6e3983a76d7/
• SHINYPROXY_USERGROUPS : DEFAULT,ZELUS,ADMIN
• SHINYPROXY_USERNAME : tan
— Installed Zelus Soccer Packages —
• fbutils (1.6.5) • fbplots (0.2.5) • fbshiny (1.0.0)
• fbmodels (0.2.52)
[...]
```

```
— Not Installed —
• fbvalidate • fbmonitoring • fbmappings
• fb scrape • qualitycontrol
[...]
```

```
— Package Options —
No options set for above packages
— Package Dependencies —
• cachem (1.0.8) • memoise (1.9.2) • timechange (0.2.0)
• data.table (1.14.8) • mgcv (1.9-0) • xgboost (1.7.5.1)
• dplyr (0.3.1) • purrr (1.0.2) • lattice (0.21-8)
• globals (0.16.2) • rlang (1.1.1) • mgcv (1.9-0)
• glue (1.6.2) • shiny (0.11.0) • nlme (3.1-162)
• jsonlite (2.2.0) • stringr (1.5.0) • splines (4.3.1)
[...]
```

[Code](#)

sitrep() |>
jsonlite::toJSON()

```
{
  "system_info": {
    "r_version": "R version 4.2.1 (2022-06-23)",
    "os_version": "Ubuntu 20.04.5 LTS"
  },
  "env_vars": {
    "ZELUS_ENV": "WEBSITE",
    "ZELUS_DB_LOCATION": "us-west-2",
    "SHINYPROXY_OIDC_ACCESS_TOKEN": "eyJ...<redacted>..z1Q",
    "SHINYPROXY_USERGROUPS": "DEFAULT,ZELUS,ADMIN",
    "SHINYPROXY_USERNAME": "tan"
  },
  "installed": [
    {"package": "fbutils", "version": "1.6.5"},
    {"package": "fbplots", "version": "1.3.0"},
    {"package": "fbmodels", "version": "0.2.52"},
    {"package": "fbshiny", "version": "1.0.0"}
  ],
  "package_options": {},
  "not_installed": [
    "fbvalidate",
    "fbmonitoring",
    "fbmappings",
    "fbscrape",
    "qualitycontrol"
  ],
  "timestamp": "2023-10-09 12:58:26"
}
```

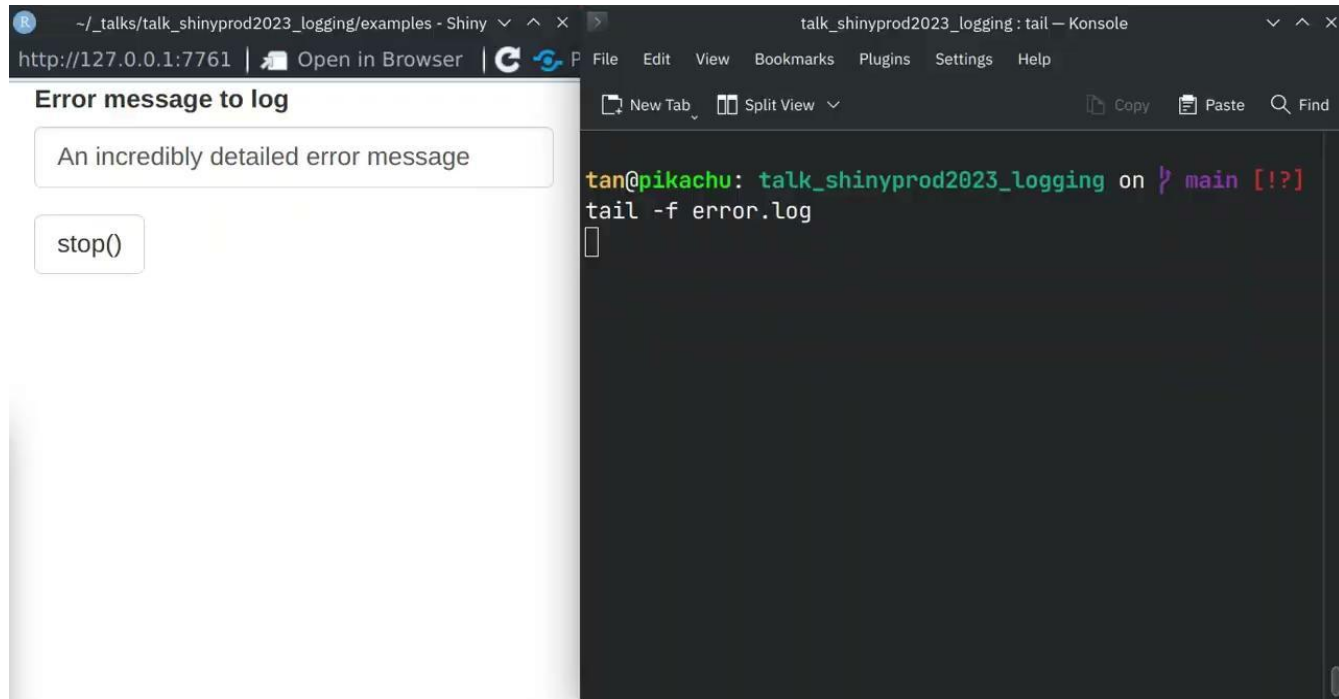
Logging app crashes as [FATAL] errors

Experimental...
hopefully not (too)
off-label? 😅

Code

```
1 log_crash ← function(){  
2   e ← get("e", envir = parent.frame())  
3  
4   stack_trace ← shiny::printStackTrace(e, full = TRUE) ▷  
5   capture.output(type = "message") ▷  
6   list()  
7  
8   logger::log_fatal(  
9     msg = e$message,  
10    stack_trace = stack_trace,  
11    timestamp = Sys.time(),  
12    sitrep = sitrep()  
13  )  
14  stop(e)  
15 }  
16  
17 options(shiny.error = log_crash)  
18
```

Logging app crashes as (FATAL) errors



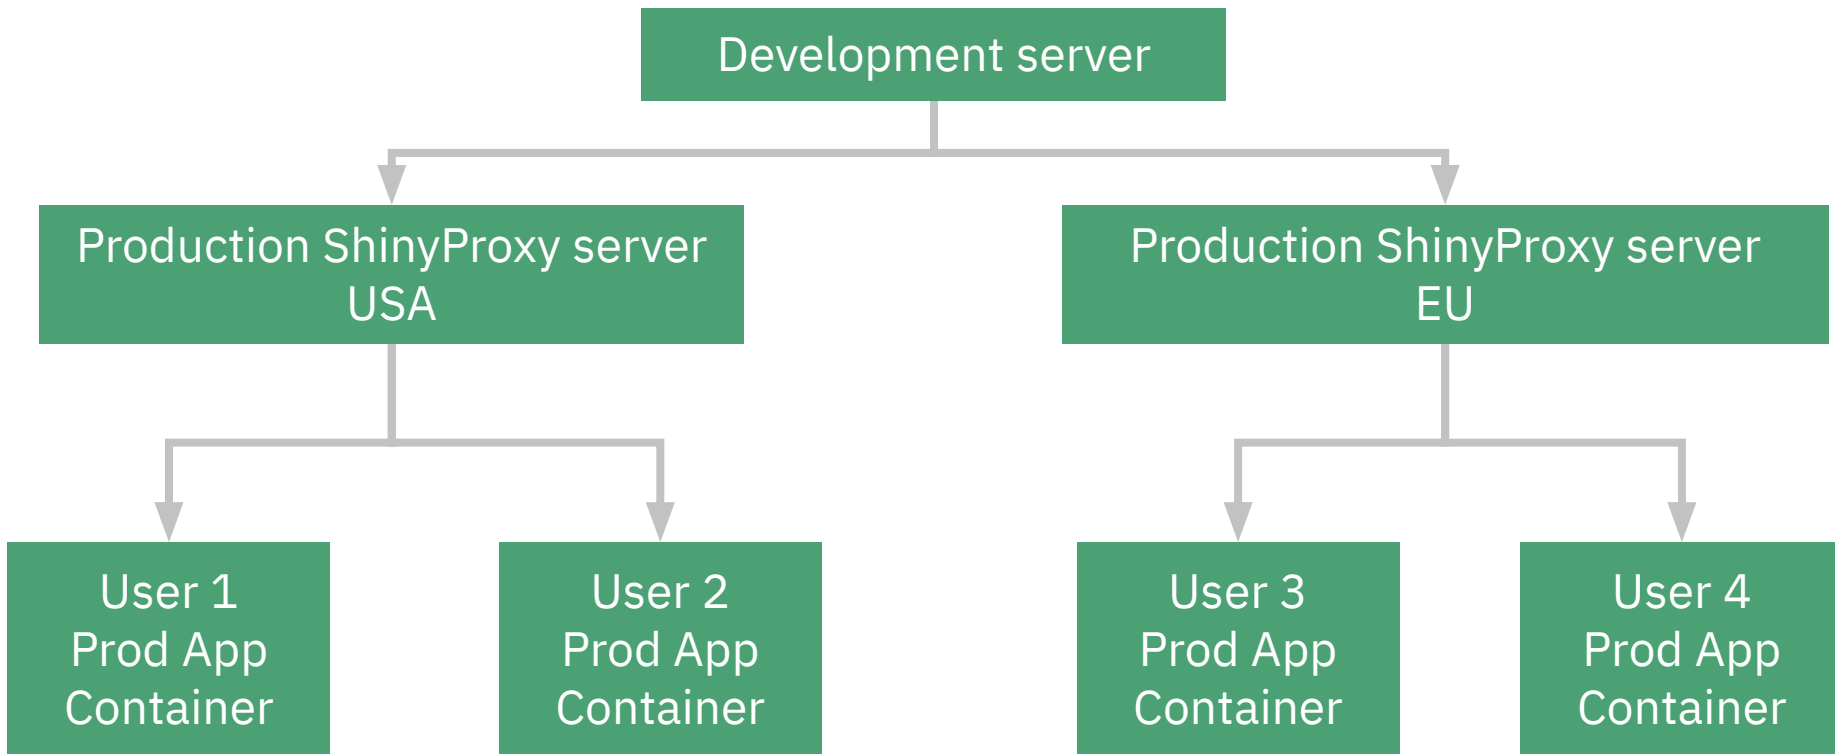
Logging as a Production System





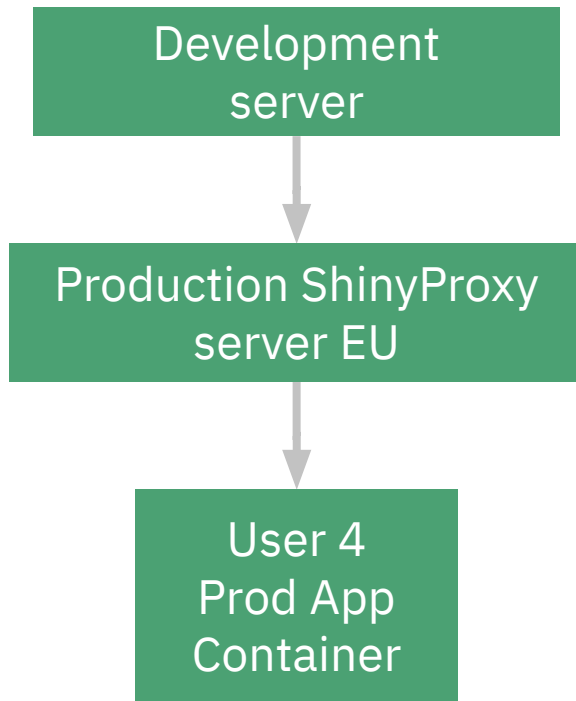
Me, trying to find the log files in our production Shiny apps

Zelus Soccer's Shiny Architecture



Digging through the logs

1. ssh into development server
2. ssh from dev into correct prod server
3. Find container's log files
4. Try to find specific error in log files
5. Pray:
 - a. that my vim-fu isn't too rusty
 - b. that I'm looking at the right set of log files
 - c. that the logs are useful



**It would be really
nice if we had a
centralized system...**

Choosing a log management system



- Grafana + Loki + Kubernetes
 - aka [ShinyProxy Monitoring Stack](#)
- AWS Cloudwatch
- Splunk
- Datadog, Sentry, Amplitude
- Other proprietary tools
- Creating our own tooling

How to choose?

~~One size fits all~~

The right solution is the one that:

- fits within the existing architecture
- has the correct tradeoff of simplicity / features
- is easy to implement
- ...actually gets used!

“It is also possible to store the container log files in an S3 back-end using the following settings...”

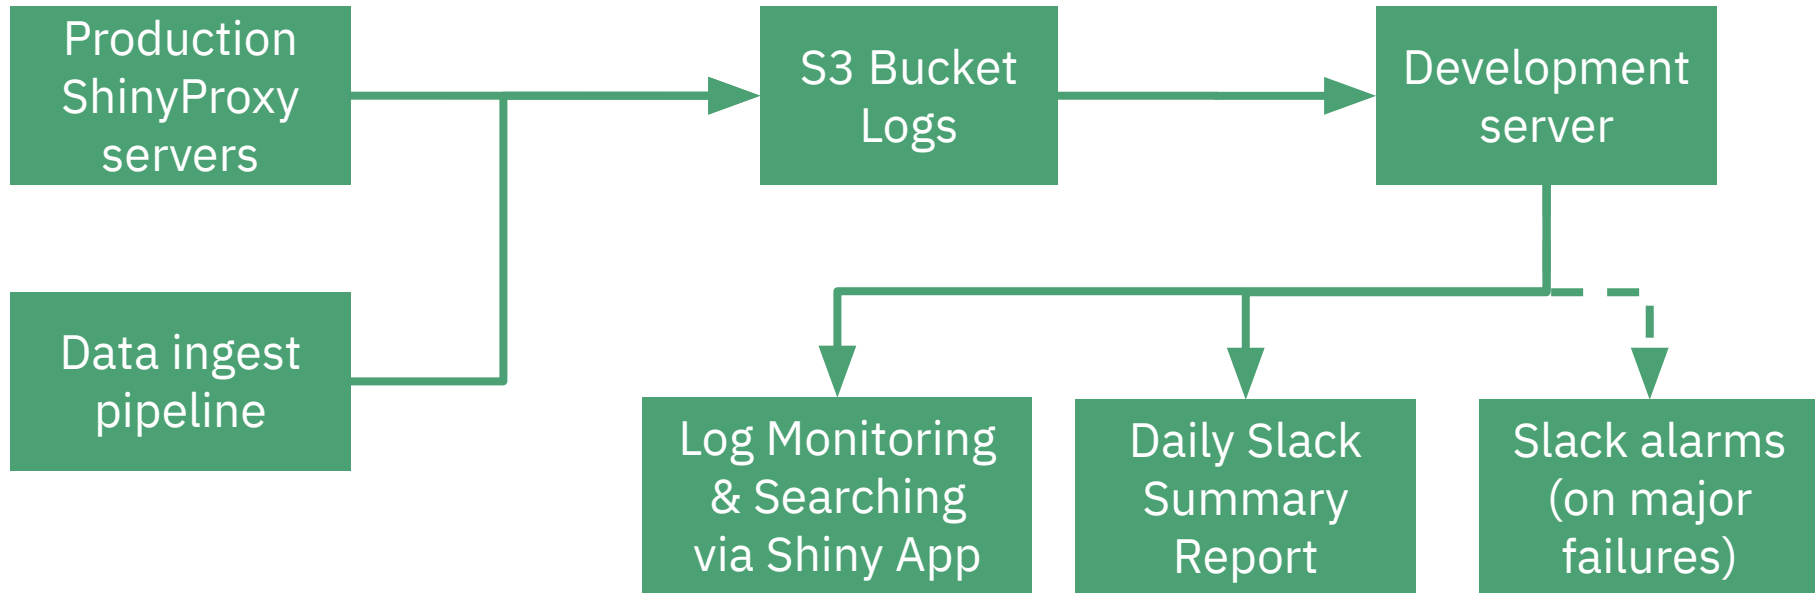


shinyproxy.io/documentation/configuration/#s3-back-end

Implementing ShinyProxy → S3 logging

ShinyProxy pushes
app logs to S3

cron job reads S3 logs into
database



How might we outgrow this system?

- Ingesting logs from other (non-R / non-JSON) systems
- Inability to use one main S3 bucket for all logs
- Log volume increases beyond capacity
- Monitoring latency/speed requirements increase
- Monitoring metrics/analysis demands increase
- ...?

Until we outgrow it, this system works for us!

Takeaways

Logging...

is about being
proactive for
future debugging

Logging...

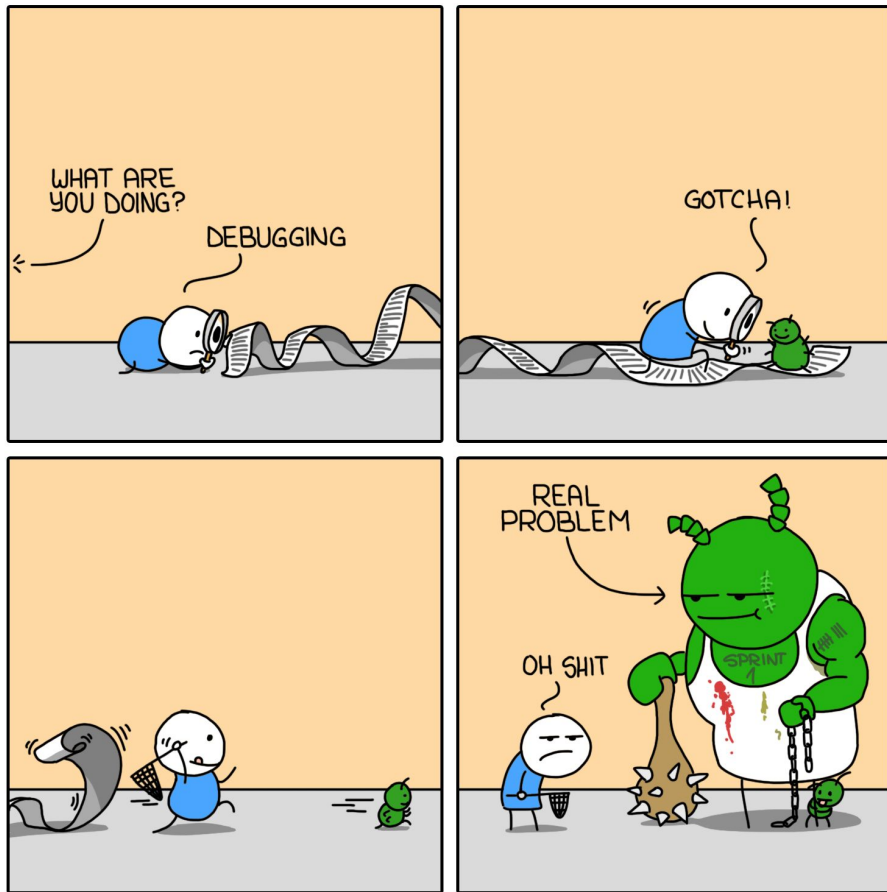
generates data
that can guide
improvements



Logging...

is a production system that you can improve on iteratively

ROOT CAUSE



MONKEYUSER.COM

<https://www.monkeyuser.com/2018/root-cause/>

Logging...

*might not always be
enough to save you from a
painful debugging
experience*

Thank you!

tanho.ca/logging-shiny
github.com/tanho63
tan@tanho.ca



Resources and Notes

R-specific tooling:

- [{logger} pkg](#) by Gergely Daróczi
- Shiny's [official docs](#) for `options(shiny.error)`
- [shinymetrics](#) by John Coene

General principles:

- [A Guide To Application Logging](#) - Andre Rabold
- [Structured logging](#) - Reflectoring
- [What Should I Log In My Application](#) - Loupe